

WhoGoesThere

Ishtiyaque Ahmad, Gwyneth Allwright, Abtin Bateni, Sabrina Tsui

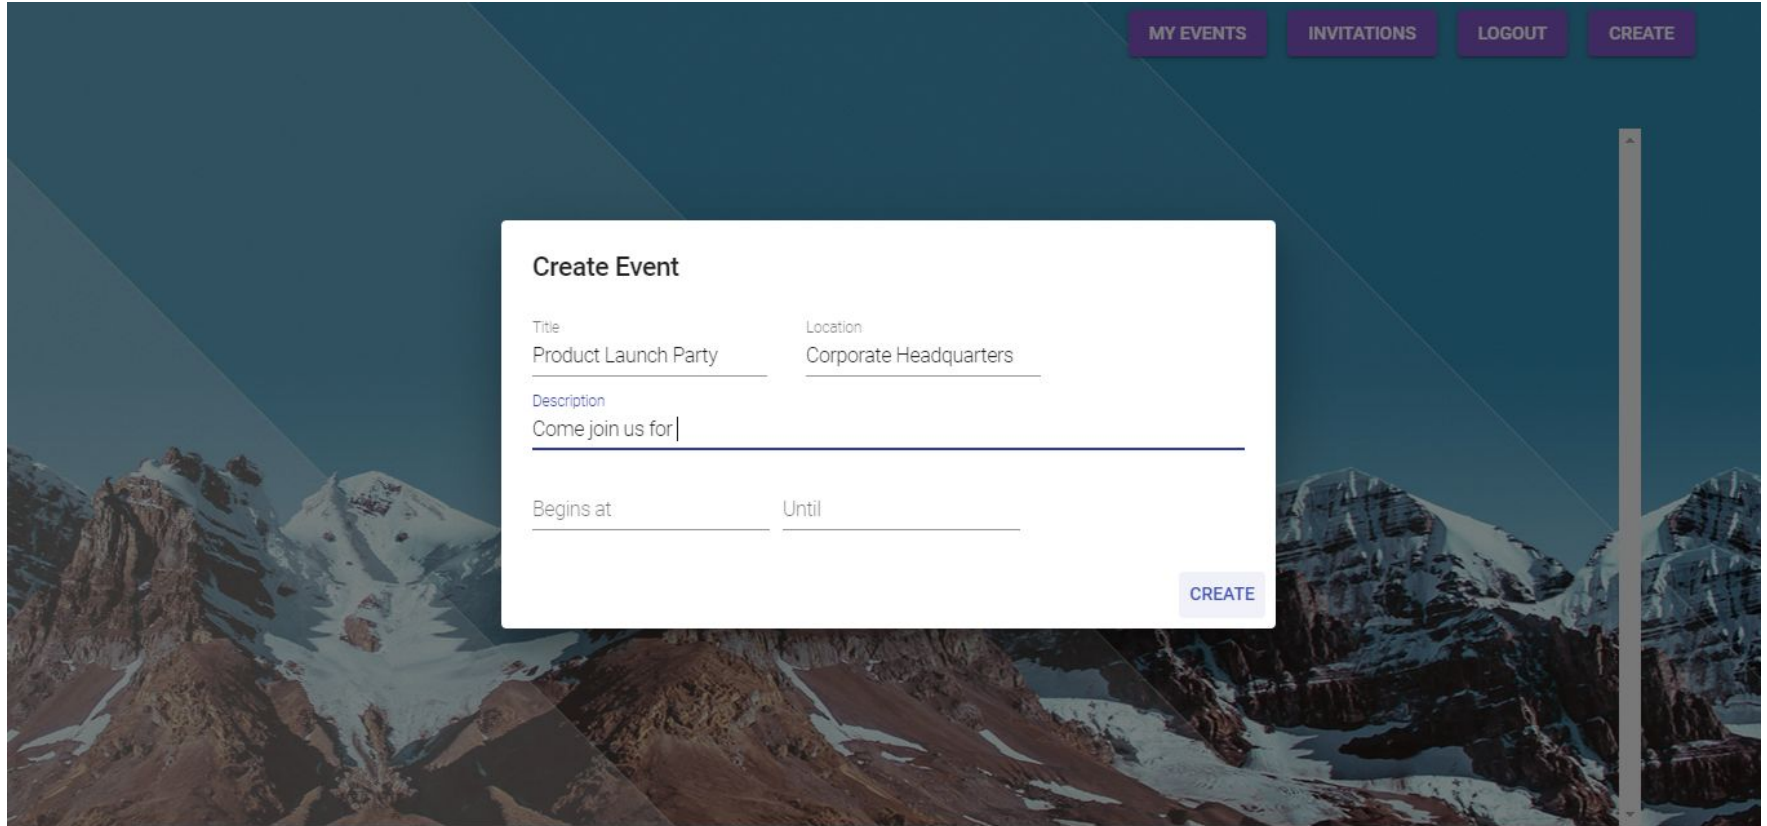
Outline

- Application Walkthrough
- Load Test Setup
- Optimizations
- Vertical Scaling Results
- Horizontal Scaling Results
- Evaluation

Events Application

- Designed for large-scale events
 - Large user base
 - Many events
- Two broad categories of users
 - Event organizers
 - Event invitees

Application Features: Creating an Event



MY EVENTS INVITATIONS LOGOUT CREATE

Create Event

Title
Product Launch Party

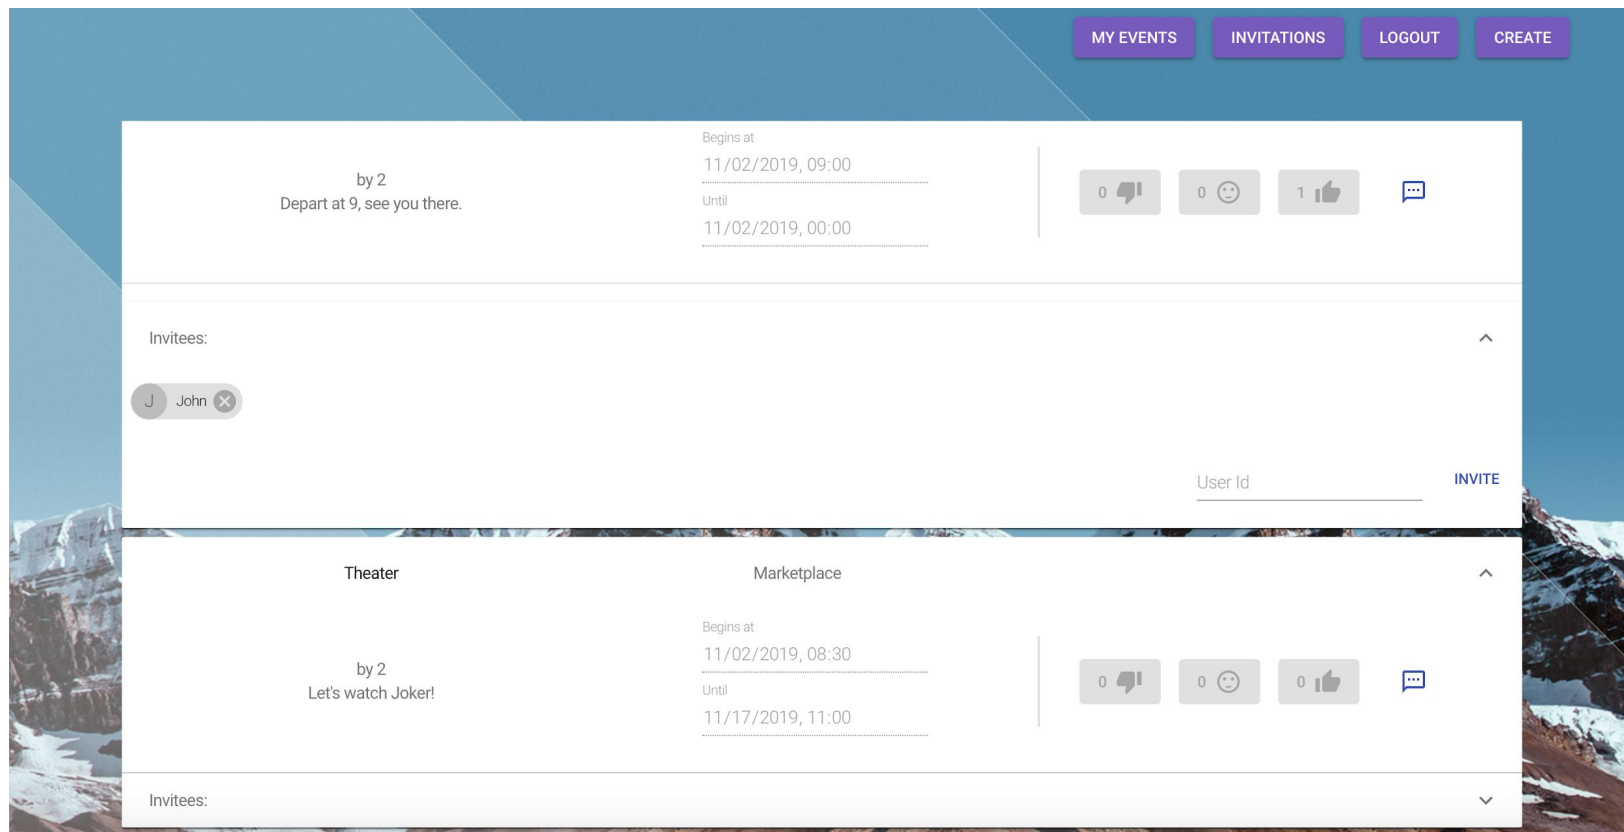
Location
Corporate Headquarters

Description
Come join us for |

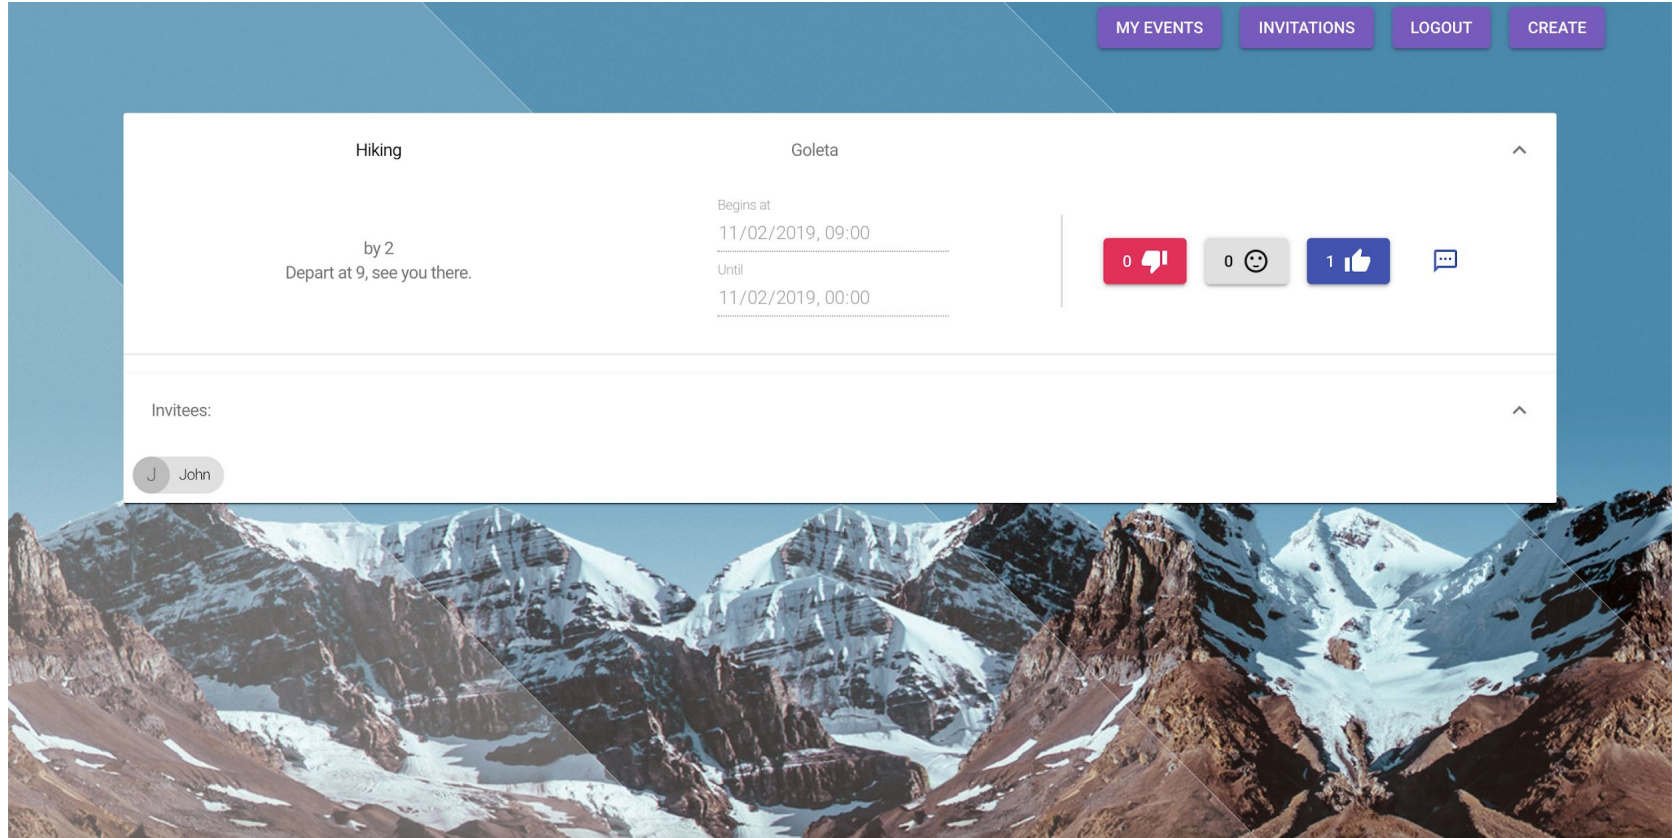
Begins at Until

CREATE

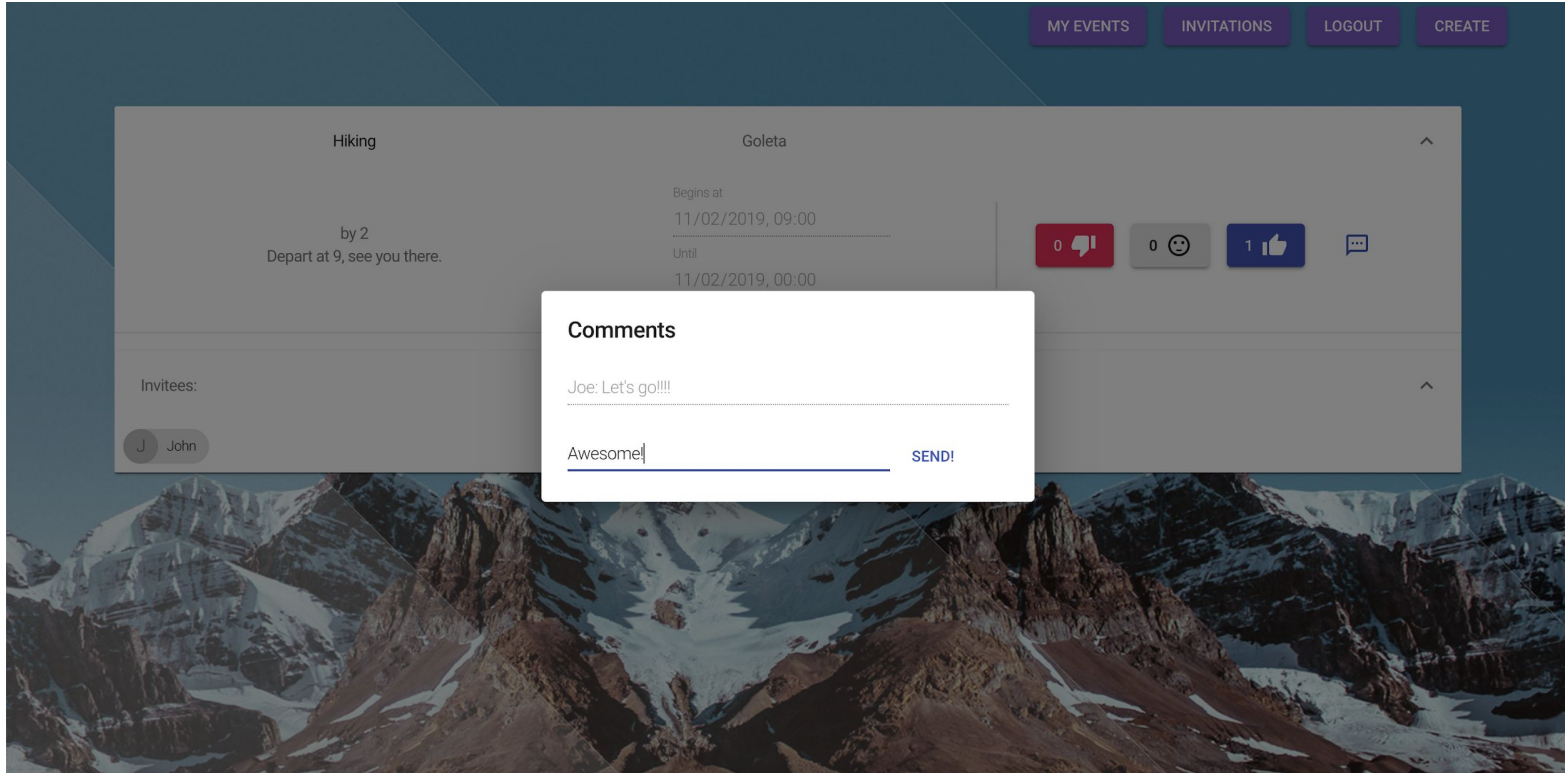
Application Features: My Events Page



Application Features: Invitations Page




Application Features: Adding a Comment



Separate Client and Backend

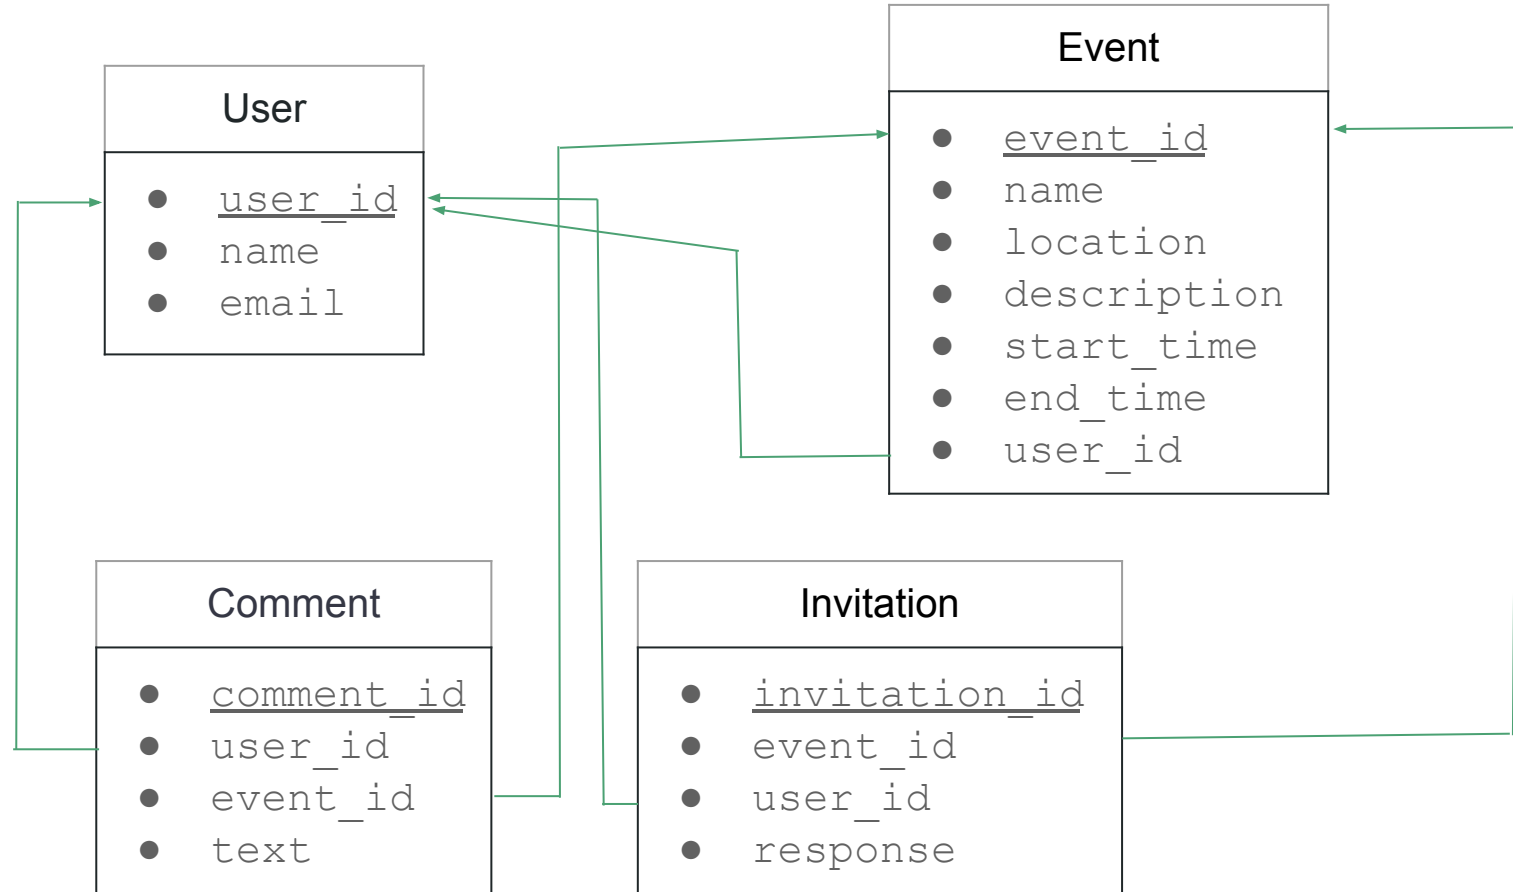
- API based service vs web pages
- Endpoints return JSON instead of html views
- Reduces code dependencies



Data Model

- Users
 - Name
 - Email
- Events
 - Name
 - Location
 - Description
 - Start time
 - End time
 - User
- Invitations
 - Event
 - User
 - Response
- Comments
 - User
 - Event
 - Text

Data Model



Tsung Test Workflow

- **Organizer (50% probability)**

- Navigate to the main page displaying all their events.
- **Create a new event.**
- Redirect to the events page with the new event.
- **Invite ten other users** to the event.

- **Invitee (50% probability)**


- Navigate to the main page displaying all their events.
- Navigate to the first event on the page.
- **RSVP to the event.**
- View the comments for the event.
- **Post a comment** to the event.

Optimizations Outline

- Database:
 - Indexing
 - Queries (JOIN versus no JOIN)
- Application Server:
 - JBuilder
 - JSON Compression (gzip)

Database Optimizations: Indexing

- Default: data is stored in a B-tree indexed by primary key
- We added indexing on non-default keys
- Faster lookup of values




Database Optimizations: Indexing

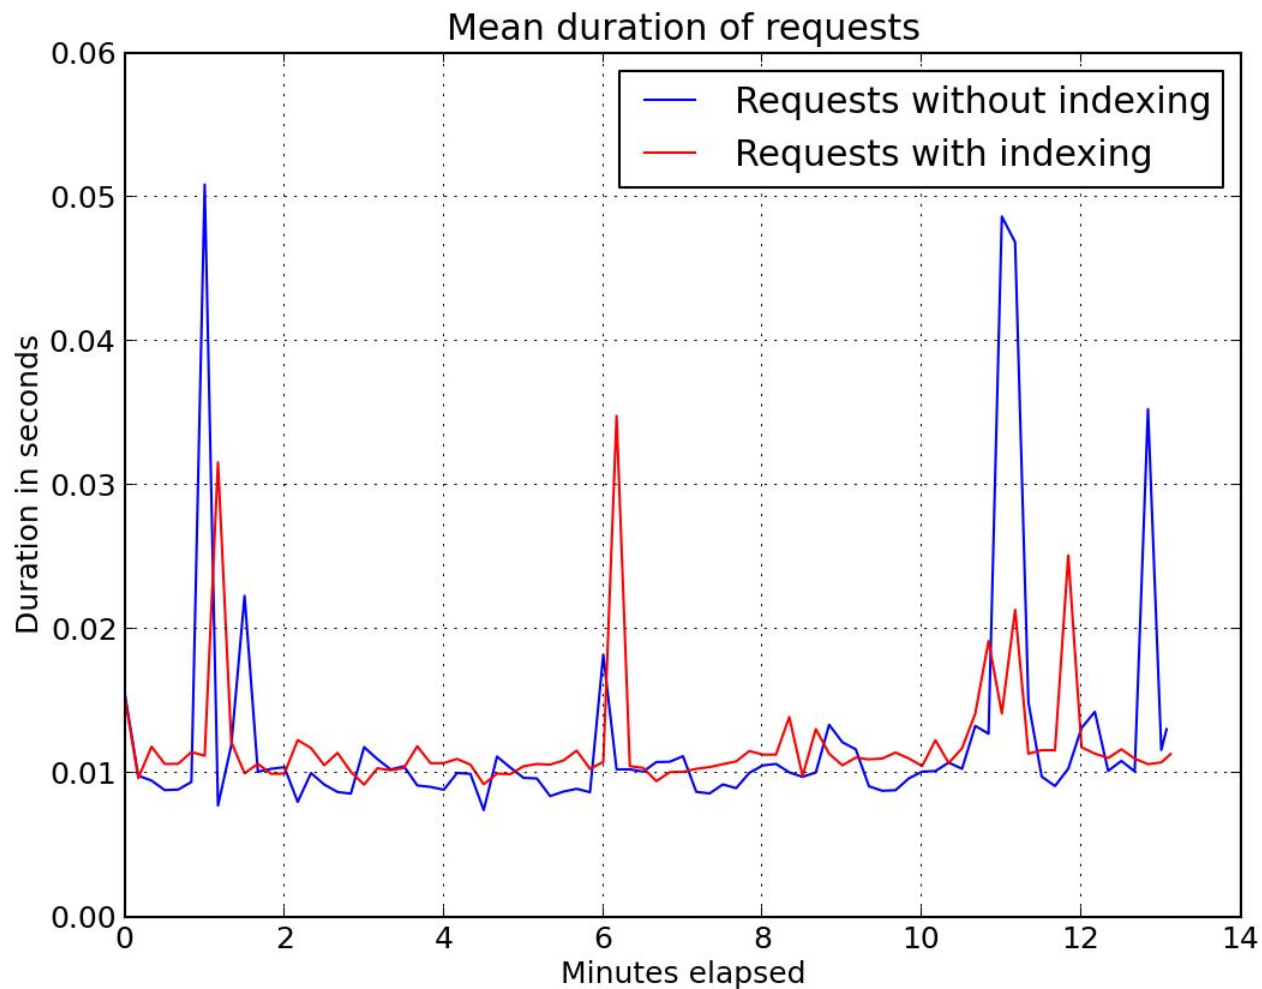
User
<ul style="list-style-type: none">● user_id● name● email

Event
<ul style="list-style-type: none">● event_id● name● location● description● start_time● end_time● user_id

Invitation
<ul style="list-style-type: none">● invitation_id● event_id● user_id● response

Comment
<ul style="list-style-type: none">● comment_id● user_id● event_id● text





Optimizations Outline

- Database:
 - Indexing
 - Queries (JOIN versus no JOIN)
- Application Server:
 - JBuilder
 - JSON Compression (gzip)

JOIN versus no JOIN

```
web_1 | ↳ app/controllers/invitations_controller.rb:9:in `indexbyuser'  
web_1 | Invitation Load (135.6ms) SELECT "invitations".* FROM "invitations"  
INNER JOIN "events" ON "events"."id" = "invitations"."event_id" WHERE "invitations"."user_id" = $1  
ns"."user_id" = $1 [{"user_id", 1}]
```

Performance with join

```
web_1 | ↳ app/controllers/invitations_controller.rb:9:in `indexbyuser'  
web_1 | Invitation Load (122.4ms) SELECT "invitations".* FROM "invitations"  
WHERE "invitations"."user_id" = $1 [{"user_id", 1}]  
web_1 | ↳ app/controllers/invitations_controller.rb:11:in `indexbyuser'  
web_1 | Event Load (27.6ms) SELECT "events".* FROM "events" WHERE "events"."  
id" IN ($1, $2, $3, $4, $5, $6, $7, $8, $9, $10, $11) [{"id", 8}, {"id", 40}, [  
"id", 25], [{"id", 38}, [{"id", 24}, [{"id", 37}, [{"id", 27}, [{"id", 4}, [{"id", 35]  
, [{"id", 50}, [{"id", 43}]
```

Performance without join

Optimizations Outline


- Database:
 - Indexing
 - Queries (JOIN versus no JOIN)
- Application Server:
 - JBuilder
 - JSON Compression (gzip)

JSON Optimization

```
Started GET "/users/" for 172.27.0.1 at 2019-11-27 17:42:29 +0000
Cannot render console from 172.27.0.1! Allowed networks: 127.0.0.0/127.255.255.255, ::1
Processing by UsersController#index as JSON
  Rendering users/index.json.jbuilder
  User Load (1.2ms)  SELECT "users".* FROM "users"
  ↳ app/views/users/index.json.jbuilder:1
  Rendered users/_user.json.jbuilder (Duration: 0.6ms | Allocations: 108)
  Rendered users/_user.json.jbuilder (Duration: 0.4ms | Allocations: 82)
  Rendered users/_user.json.jbuilder (Duration: 0.4ms | Allocations: 82)
  Rendered users/_user.json.jbuilder (Duration: 0.7ms | Allocations: 82)
  Rendered users/_user.json.jbuilder (Duration: 1.2ms | Allocations: 82)
  Rendered users/_user.json.jbuilder (Duration: 4.3ms | Allocations: 82)
  Rendered users/_user.json.jbuilder (Duration: 0.7ms | Allocations: 82)
  Rendered users/_user.json.jbuilder (Duration: 0.6ms | Allocations: 82)
  Rendered users/_user.json.jbuilder (Duration: 0.5ms | Allocations: 82)
  Rendered users/_user.json.jbuilder (Duration: 0.5ms | Allocations: 82)
  Rendered users/index.json.jbuilder (Duration: 40.5ms | Allocations: 3406)
Completed 200 OK in 46ms | Views: 41.4ms | ActiveRecord: 1.2ms | Allocations: 3910)
```


JSON Optimization

- JBuilder Optimization in Rails



JSON Optimization

```
app/controllers/users_controller.rb:8:in `index'
Started GET "/users/" for 172.27.0.1 at 2019-11-27 17:59:49 +0000
Cannot render console from 172.27.0.1! Allowed networks: 127.0.0.0/127.255.255.255, ::1
Processing by UsersController#index as JSON
  Rendering users/index.json.jbuilder
  User Load (0.5ms)  SELECT "users".* FROM "users"
    ↳ app/views/users/index.json.jbuilder:1
  Rendered users/index.json.jbuilder (Duration: 29.3ms | Allocations: 5136)
Completed 200 OK in 35ms  Views: 23.0ms  ActiveRecord: 8.1ms | Allocations: 8853)
```


Optimizations Outline

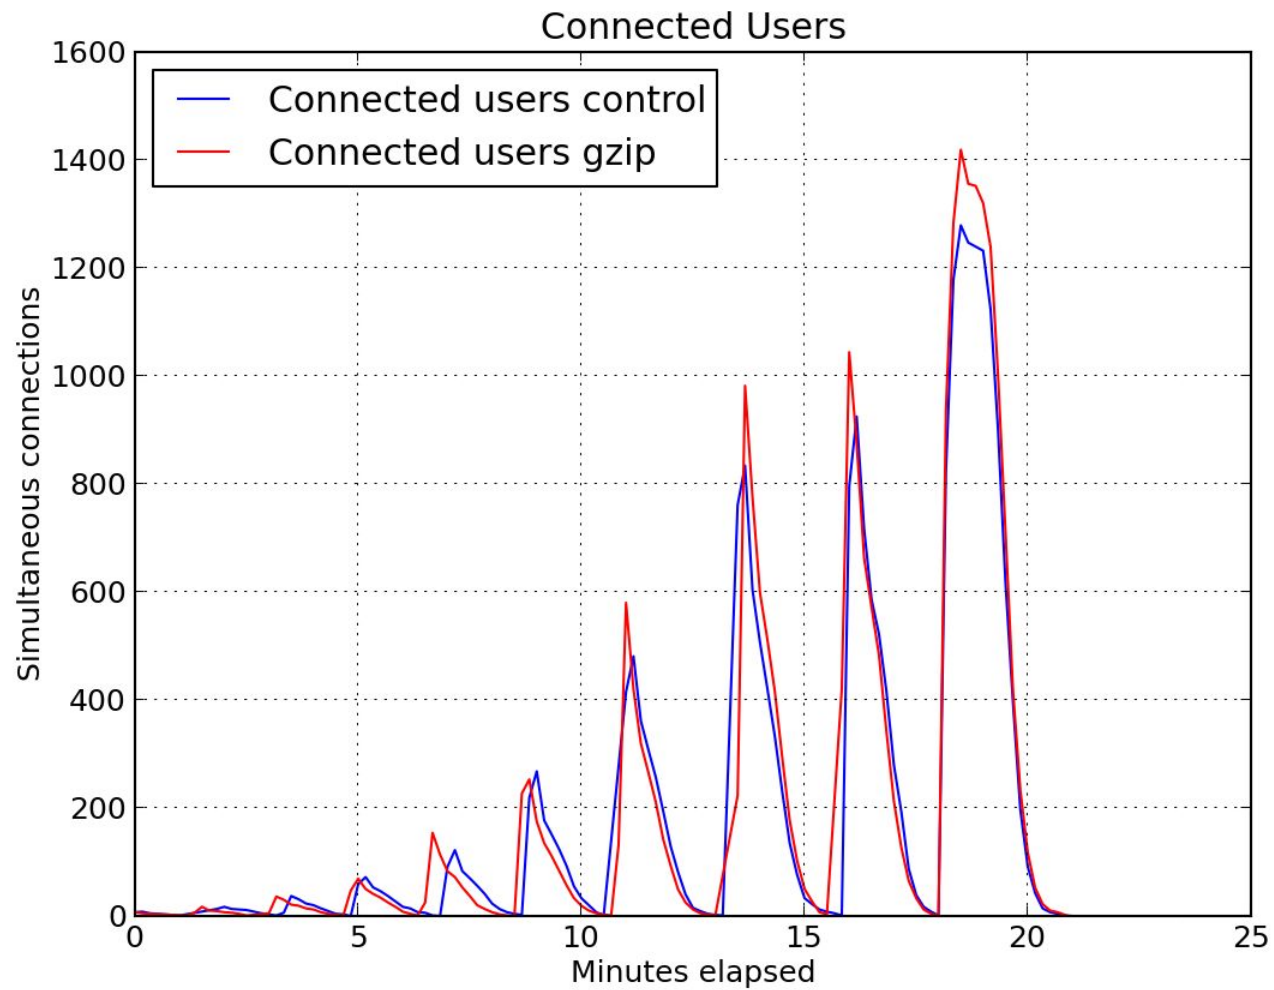
- Database:
 - Indexing
 - Queries (JOIN versus no JOIN)
- Application Server:
 - JBuilder
 - JSON Compression (gzip)

JSON Compression

```
GET /encrypted-area HTTP/1.1  
Host: www.example.com  
Accept-Encoding: gzip, deflate
```


```
HTTP/1.1 200 OK  
Date: mon, 26 Jun 2016 22:38:34 GMT  
Server: Apache/1.3.3.7 (Unix) (Red-Hat/Linux)  
Last-Modified: Wed, 08 Jan 2003 23:11:55 GMT  
Accept-Ranges: bytes  
Content-Length: 438  
Connection: close  
Content-Type: text/html; charset=UTF-8  
Content-Encoding: gzip
```





Vertical Scaling Results

Application Instance	Database Instance	Max Arrival Rate Before 5XX Error (users/second)
1x m5.large	1 x db.m5.large	100
	1 x db.m5.xlarge	110
	1 x db.m5.4xlarge	120
1x m5.4xlarge	1 x db.m5.large	610
	1 x db.m5.xlarge	620
	1 x db.m5.4xlarge	620
1x m5.8xlarge	1 x db.m5.large	1280
	1 x db.m5.xlarge	1350
	1 x db.m5.4xlarge	1760

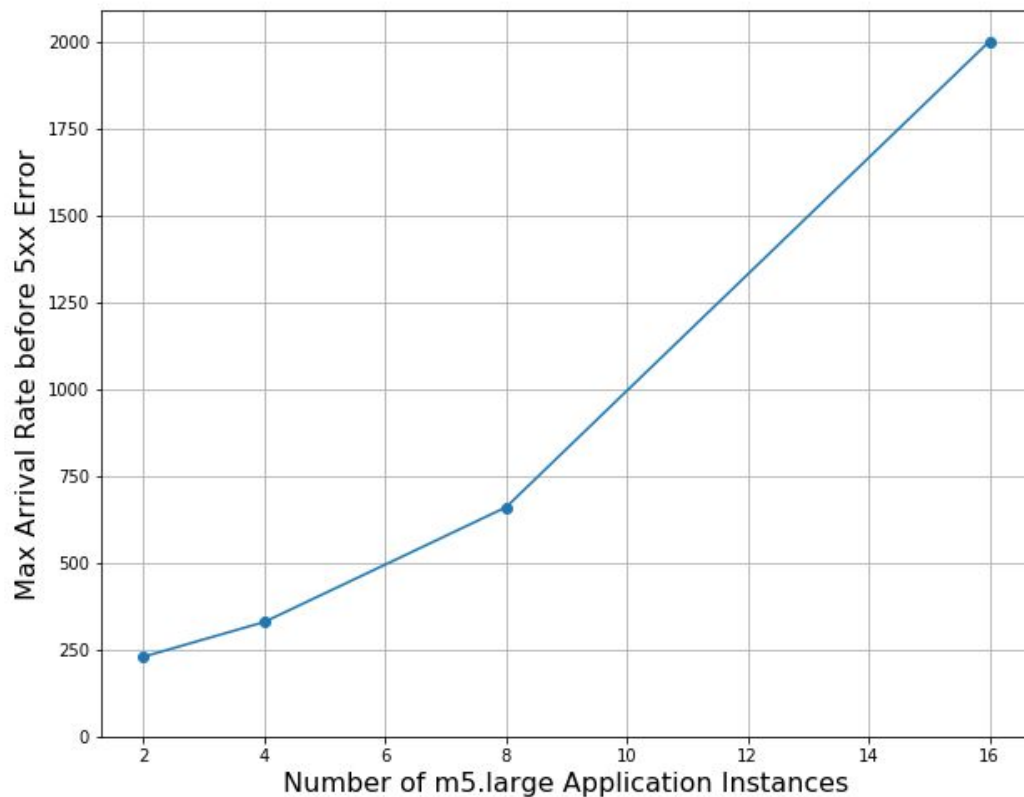


Horizontal Scaling Results

Number of m5.large Application Instances	Max Arrival Rate Before 5XX Error	Max Users/s per Dollar
2	230	622
4	330	587
8	660	698
16	>2000	>1167

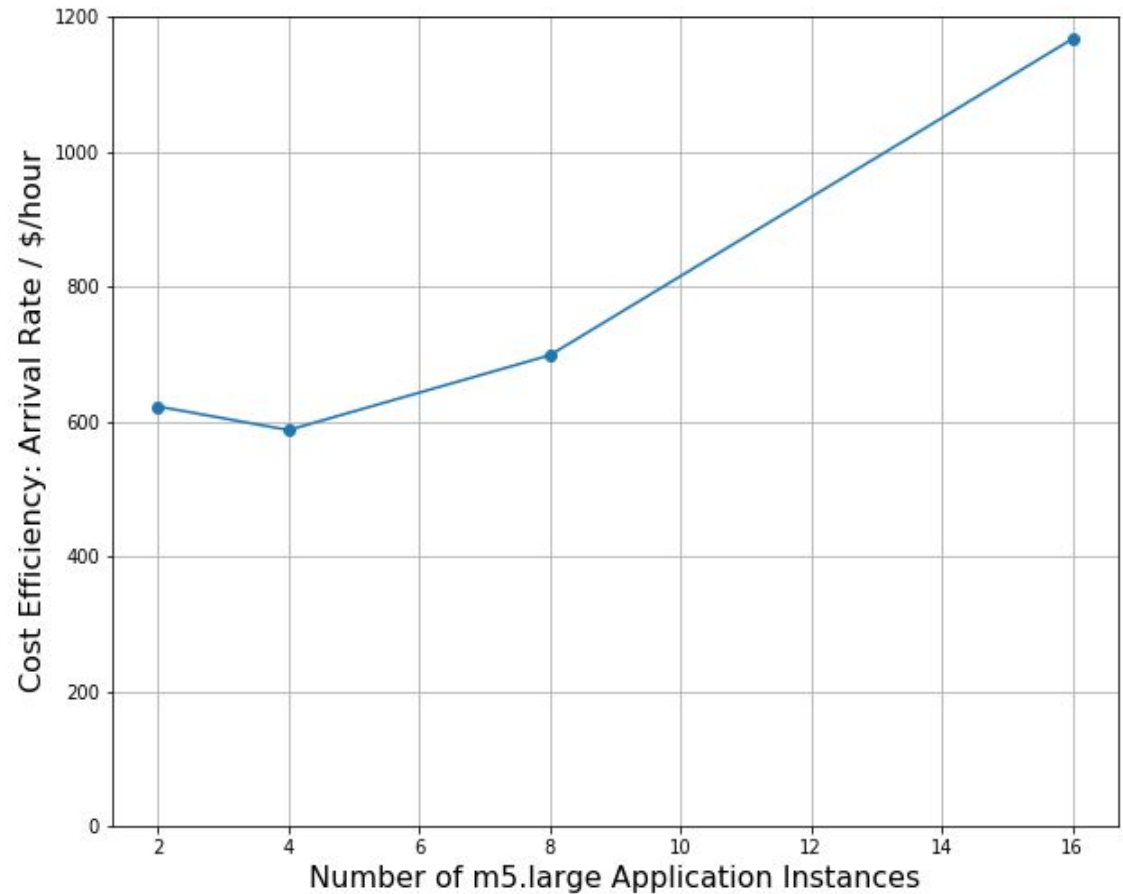
Horizontal Scaling Results

- Impact on traffic



Horizontal Scaling Results

- Impact on cost



Evaluation

- Combining all optimizations and scaling the final application
- Baseline performance
 - m5.large instance
 - Max 90 users/second tsung arrival rate
- Optimized performance
 - Adding optimizations: improvement of 10 users/second over baseline
 - Best performance: 16 x m5.large instance + db.m5.large
 - Max >2000 users tsung arrival rate

Thank You

References

- https://en.wikipedia.org/wiki/HTTP_compression
- <https://commons.wikimedia.org/wiki/File:B-tree-definition.png>